


Tutorial 3

A magestic introduction to the 3rd assignment

Topics of 3rd assignment





Transfer Learning - Important Q&As

- **What is Transfer Learning?**

- Transfer learning is a machine learning technique where a model trained on one task is re-purposed on a second related task.

- **Can we do this for every dataset?**

- No, the original dataset (in which the NN has been trained) and the one in which we want to apply it **MUST** have similar domains

- **Do we use the NN architecture as it is?**

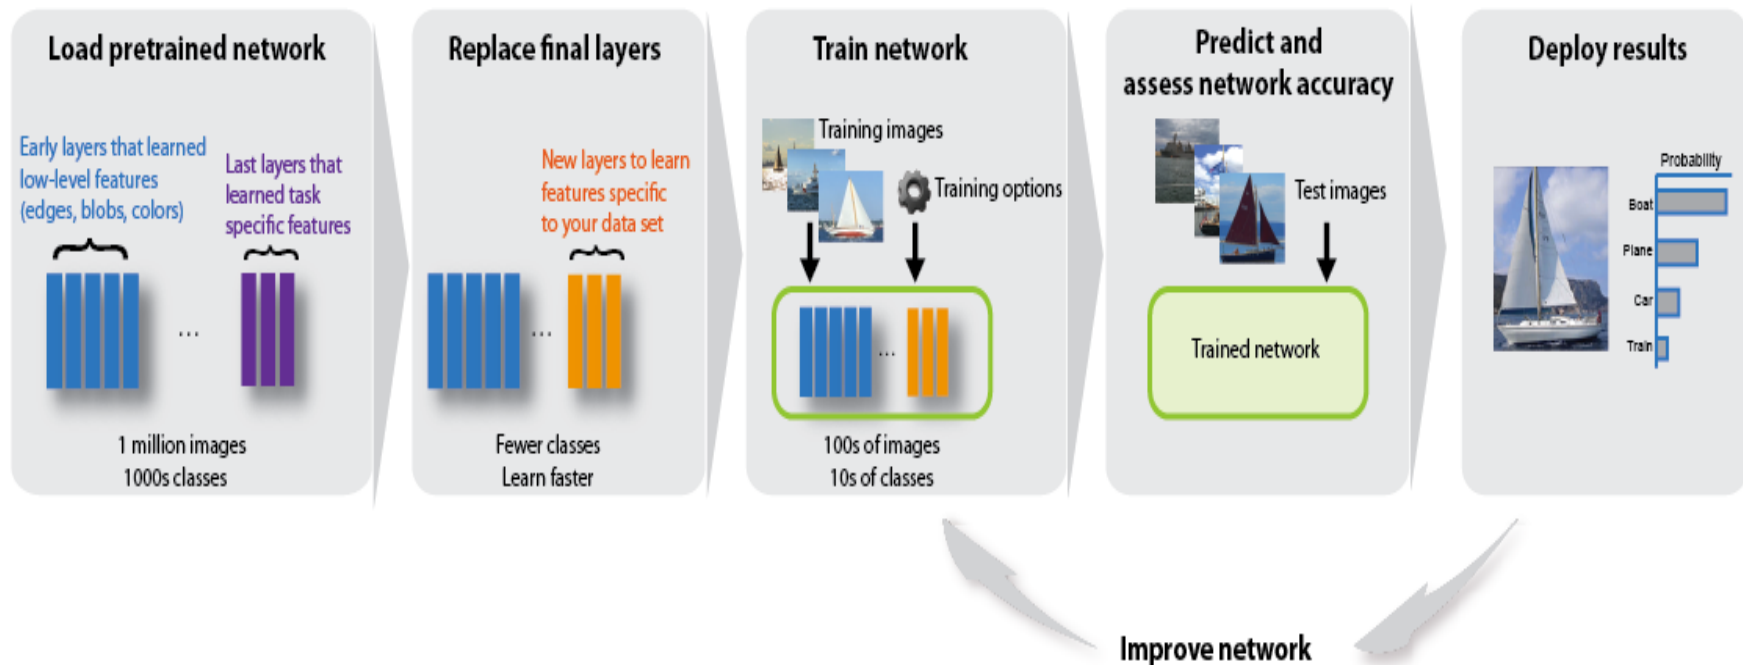
- No, early layers are more generic and later layers are original-dataset-specific !
- We keep the lower layers! Toss the higher ones!

- **So, how we do it?**

- Load the model pretrained on the original dataset (usually on ImageNet)
- Remove higher layers (usually the FC ones)
- Add our dataset-specific layers on top
- Freeze learning for the lower layers (original) and training the new additions on the new dataset!

Transfer Learning – Basic Concepts

Reuse Pretrained Network





More Q&As - *we will answer them together!*

- **My dataset is small but is similar to the original dataset. Should I fine-tune?**
 - *Answer:*
 - And how about my hyperparameters?
- **My dataset is large and similar to the original dataset. Should I fine-tune or train from scratch?**
 - *Answer:*
- **My dataset is different from the original. Should I finetune?**
 - *Answer:* *(Hint: if it is small we can do something)*
- **How do I decide what to put on - top?**
 - *Answer:*



Task 3.1: Fine-tune a model

- Which model? AlexNet!
- Layers and specs:
 - **Layer 1:**
 - Input: $224 \times 224 \times 3$
 - Number of filters: 96
 - Kernel: $11 \times 11 \times 3$, Stride 4
 - Output: $55 \times 55 \times 96$, Act: ReLU
 - **Layer 2:**
 - Input: $55 \times 55 \times 96$
 - Max - pooling (stride 2) -> Conv
 - Number of filters: 256
 - Kernel: $5 \times 5 \times 48$
 - Output: $27 \times 27 \times 256$
 - **Layers 3, 4 and 5 are similar to 2!**
 - **Layer 6:**
 - Fully Connected: Input -> $13 \times 13 \times 128$
 - Output: 1×2048




Image Source:

<http://www.cs.toronto.edu/~fritz/absps/imagenet.pdf>

- **Layers 7 and 8 are similar with 8** outputting $1 \times N_{\text{classes}}$

Task 3.1: Dataset Description

- WikiArt Dataset
 - 4000 images of paintings
 - 10 different styles (i.e. Baroque, Realism, Expressionism)



Task 3.1: Steps to be followed

- Download the WikiArt dataset (download_wikiart.py file)
 - Be patient, it may take a while...
- Download AlexNet, with the specific weights learned by the network
- Implement the loss function, the optimization process with GD algorithm, and the accuracy obtained
 - Compute the gradients of all variables, and apply GD algorithm
- Fine-tune AlexNet network, by subtracting:
 - The last Fully-Connected (FC) layer
 - The last two final FC layers
- Report accuracy, loss (training & validation) ++ intuition gained
- Epochs: 10 → Training may take a while, gradual improvement though

Image Class Saliency Map Visualization

- Saliency Map: Measures the degree to which each pixel in the image affects the classification score for that image



Image Source:
<https://arxiv.org/pdf/1312.6034.pdf>.

Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps", ICLR Workshop 2014.

Image Class Saliency Map Visualization

- How do we compute it?
 - Compute the gradient w.r.t image pixels of the un-normalized score corresponding to the correct class
 - Image : $(H,W,3) \rightarrow$ Gradient : $(H,W,3)$
 - Obtain absolute value of each of these gradients, and take the maximum value over 3 possible image channels (RGB)

Image Class Saliency Map Visualization

- For every image:

- Do a forward pass of the image through the network

$$S_c(I) = w_c^T I + b_c$$

- Calculate the scores for every class

- Enforce the derivative w.r.t. I (i.e. input image) of score vector S at last layer for all classes except class C to be 0, while for C set it to 1

- Back-propagate this derivative until the start $\arg \max_I S_c(I) - \lambda \|I\|_2^2 \quad w = \frac{\partial S_c}{\partial I} \Big|_{I_0}$

- Render the gradients and obtain the Saliency Map

- Network used: VGG_16 (trained on a subset of the ImageNet database)